

Project 1

15-441: Computer Networks

Matt Mukerjee
David Naylor
Ben Wasserman

Agenda

- Handling Concurrency
- Project 1 Checkpoint 1
- Q & A

Flashback!

- `getaddrinfo()` - Prepare to launch!
- `socket()` - Get the file descriptor!
- `bind()` - Which port am I on?
- `listen()` - Will someone please call me?
- `connect()` - Hey, you!
- `accept()` - Thank you for calling port 8080!
- `send()` and `recv()` - Talk to me, please!
- `close` and `shutdown()` - Get out!

What do you want to build?

A webserver that can handle multiple concurrent connections!

What's the problem?
Blocking!

What's the solution?
Threading or `select()`

Threading approach

- Did in 15-213??
- Main server blocks on `accept()`
- Accept incoming connection
- `Fork()` child process for each connection
- **Pain!**
 - Need to manage a pool of threads
 - And what if tasks have to communicate?

World of `select()`

- Event driven programming!
- Single process that `multiplexes` all requests.
- Caveat
 - Programming is not so transparent!
 - Server no longer acts like it has only one client!

How to use `select()`?

- Give select a set of sockets/file descriptors.
- `select()` blocks till `something` happens.
 - Data coming in on some socket.
 - Able to write to a socket.
 - Exception at the socket.
- Once woken up, check for the event and `service` it the way the server would do.

select()

```
#include <sys/select.h>
```

```
int select (int nfd, fd_set* readfds,  
            fd_set* writefds, fd_set* exceptfds,  
            struct timeval *timeout);
```

fd_set Datastructure

- Remember, file descriptor is just an integer!
- Datastructure is basically a bit array!
- Helper macros:

`FD_ZERO(fd_set* fdset);` /* initializes fdset to have 0s for all fds */

`FD_SET(int fd, fd_set* fdset);` /* sets the bit for fd in fdset */

`FD_CLR(int fd, fd_set* fdset);` /* clears the bit for fd in fdset */

`FD_ISSET(int fd, fd_set* fdset);` /* returns non-0 if fd is set else 0 */

select() Parameters

- The FDs between 0 to `nfds`-1 are checked.
- Check for reading in `readfds`.
- Check for writing in `writefds`.
- Check for exception in `exceptfds`.
- These `fd_sets` can be `NULL`.
- `timeout`
 - `NULL` – blocking
 - else how long to wait for the required condition before returning to the caller.

Return value, Error states

- Success – number of ready descriptors.
 - `readfds`, `writefds` and `exceptfds` are modified
- Time expired – returns 0 (errno set to `EINTR`)
- Failure – returns -1
 - `EBADF`, `EINTR`, `EINVAL`, `ENOMEM`

Pseudo-code of Usage

- `nfds` = 0
- Initialize `readfds`, `writefds`, `exceptfds` using `FD_ZERO`
- Add the listener socket to `readfds` using `FD_SET` and update `nfds`
- For each active connection
 - If connection has available read buffer, add fd to `readfds` (`FD_SET`)
 - If connection has available write buffer, add to `writefds` (`FD_SET`)
 - Add to `exceptfds` (`FD_SET`) – not really needed for this project.
 - Update `nfds` to ensure that the fd falls in the range
- `select_return` = `select(nfds, readfds, writefds, exceptfds, NULL)`
- If `select_return` > 0
 - Handle exceptions if any fd in `exceptfds` is set to 1 (`FD_ISSET`)
 - Read data from connections for which fd in `readfds` is set to 1 (`FD_ISSET`)
 - Write data from connections for which fd in `writefds` is set to 1 (`FD_ISSET`)
 - If listener socket is set to read, `accept` and handle new connection.
- Else handle error states

cp1_checker.py

- `./cp1_checker.py` <ip> <port> <#trials> <#writes and reads per trial> <max # bytes to write at a time> <#connections>
 - Starts #connections connections to server at ip and port
 - Repeat #trials number of times
 - Sample #writes and reads per trials connections.
 - Send random number of random bytes to each of these connections (with a limit of max # bytes to write at a time).
 - Receive and check if all the bytes received are same as the ones that are sent.
 - If your server cannot handle multiple connections
 - Set #connections to 1 and #writes and reads per trial to 1

Okay, so you can handle multiple connections!
But that is not enough...

Reading data

- Check return value of `recv()`
 - Error – handle the error and clear up state.
 - If peer shutdown the connection, clear up state.
- Maintain state
 - Maintain a read buffer
 - Keep track of the number of bytes left to be read
 - May need multiple reads to get all data
 - But only one read per socket when `select()` returns.

Writing data

- Check return value of `send()`
 - Error – handle the error and clear up state.
 - If peer shutdown the connection, clear up state.
- Maintain state
 - Maintain a write buffer
 - Keep track of the number of bytes left to be written
 - May need multiple writes to send all data
 - Number of bytes actually sent should be checked from the return value
 - Only one write per socket when `select()` returns.

Exceptfds

- For handling out of band data
- Should be read one byte at a time!
- Not really needed for this project.

Checkpoint 1 Docs

- **Makefile** - make sure nothing is hard coded specific to your user; should build a file which runs the echo server (name it lisod)
- **All of your source code** - all .c and .h files
- **readme.txt** - file containing a brief description of your current implementation of server
- **tests.txt** - file containing a brief description of your testing methods for server
- **replay.test** - a file containing bytes that can be sent to your server as a test case
- **replay.out** - a file containing expected bytes that should be sent as a response from your server when provided replay.test
- **vulnerabilities.txt** - identify at least one vulnerability in your current implementation

Remember

- Code quality
- Code documentation
- Robustness
 - Handle all errors
 - Buffer overflows
 - Connection reset by peer

Peek into the future

- Checkpoint 2
 - Implement HTTP 1.1 parser and persistent connections
- Checkpoint 3
 - Implement HTTPS handshaking and persistent connections via TLS
 - Implement CGI server-side.

Homework 1

- Posted today
- Due September 19th
 - (A week from next Thursday)
- **Hard copy** in class **only**
 - (Will not accept emailed submissions)

Any other questions?

Come to our office hours!

Matt Gates 7501 Wednesday 11-Noon

David Gates 7509 Tuesday 4:30 – 5:30 PM

Ben Gates 4126 Thursday 3 – 4 PM